

Optimizing Facial Emotion Recognition

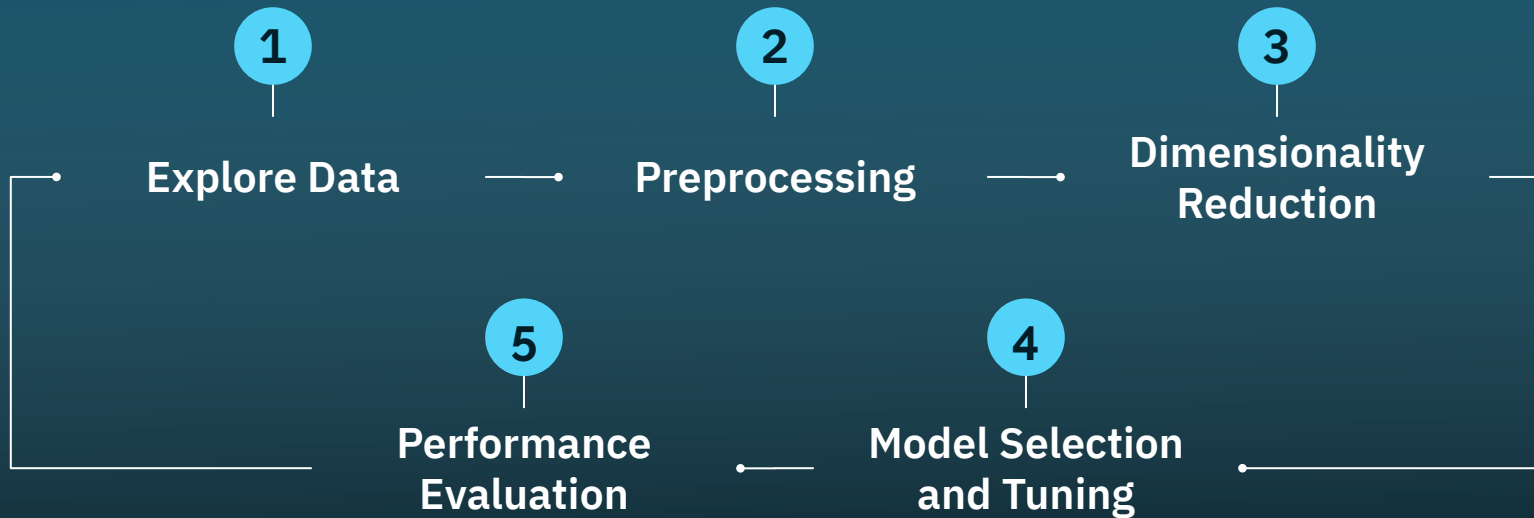
Edmond Anderson, Suliah Apatira, Jasmine Harris
CMSE 831: Computational Optimization
Longxiu Huang

Facial Image Classification

Facial emotion image classification is an AI-driven field focused on automatically recognizing and categorizing emotions expressed in human faces using advanced algorithms and deep learning models.

This technology has the potential to revolutionize human-computer interaction, mental health applications, and our understanding of emotions across various contexts.

Strategy





01 Exploratory Data Analysis

About the Dataset

- **6 Expressions:** Surprise, Anger, Fear, Happy, Neutral, Sad
- **Dimensions:** 48 x 48 pixel grayscale images

35,340 Images

Train Set: 28,273

Test Set: 7,067

Sample Images

Angry Face Neutral Face Neutral Face Surprise Face Neutral Face



Angry Face Happy Face Angry Face Angry Face Sad Face




Suprise Face Sad Face Happy Face Fear Face Suprise Face




Emotion Breakdown

80% Train



20% Test





02

Preprocessing

Preprocessing

- 30% augmented data added
- Height randomly shifted by up to 20%
- Width randomly shifted by up to 20%
- Images randomly flipped horizontally
- Images randomly rotated left/right by up to 20 degrees
- RGB values normalized (0-255) -> (0-1)

Angry Face Surprise Face Neutral Face Fear Face Happy Face



Sad Face



Fear Face



Fear Face



Surprise Face Neutral Face



Sad Face



Happy Face Neutral Face Happy Face



```
graph LR; A["Train Set:  
(28273, 48, 48)"] --> B["Augmented Set:  
(36754, 48, 48)"]; B --> C["Scaled Set  
(36754, 2304)"]
```

Train Set:
(28273, 48, 48)

Augmented Set:
(36754, 48, 48)

Scaled Set
(36754, 2304)



04

Dimensionality Reduction

Dimensionality Reduction

Principal Component Analysis

A mathematical technique to reduce the dimensionality of data. It works on the principal of factoring matrices to extract the principal pattern of a linear system

Linear Discriminant Analysis

This classical supervised dimensionality reduction method seeks an optimal linear transformation that maps the original data to a low-dimensional space.


```
graph LR; A["Scaled Set:  
(36,754, 2304)"] --> B["After PCA:  
(36,754, 856)"]; B --> C["After LDA  
(36754, 5)"]
```

Scaled Set:
(36,754, 2304)


After PCA:
(36,754, 856)

After LDA
(36754, 5)


PCA



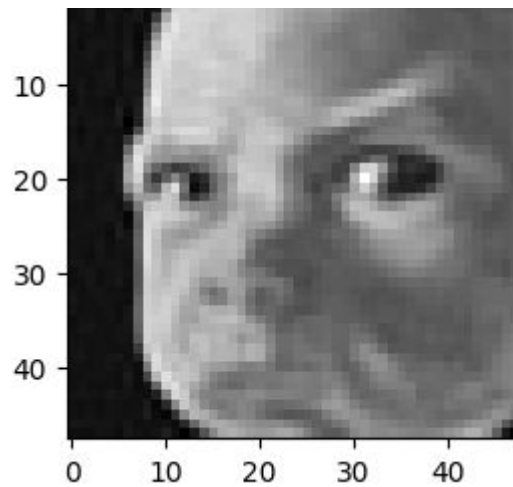


PCA



Original



Transformed



LDA

Between Class Variance

$$S_b = \sum_{i=1}^g N_i (\bar{x}_i - \bar{x})(\bar{x}_i - \bar{x})^T$$

Within Class Variance

$$S_w = \sum_{i=1}^g (N_i - 1) S_i = \sum_{i=1}^g \sum_{j=1}^{N_i} (x_{i,j} - \bar{x}_i)(x_{i,j} - \bar{x}_i)^T$$

Lower Dimensional Space

$$P_{lda} = \arg \max_P \frac{|P^T S_b P|}{|P^T S_w P|}$$



05

Model Selection and Tuning

Stochastic Gradient Descent

Model Overview

$$Q(w) = \frac{1}{n} \sum_{i=1}^n Q_i(w),$$

Pros:

- Efficiency and speed
- Convergence rate
- Regularization
- Memory efficiency

Cons:

- Noisy updates
- Varied convergence
- Learning rate running
- Sensitivity to initial conditions

Hyperparameters Tested:

- **'alpha'**: [0.0001, **0.001**, 0.01],
- **'loss'**: ['hinge', **'log_loss'**, 'modified_huber'],
- **'penalty'**: ['l1', **'l2'**, 'elasticnet'],
- **'max_iter'**: [1000, 2000, **3000**],
- **'tol'**: [1e-3, 1e-4, **1e-5**],

Best Model:

Validation Score: 0.3795749497731347

Test Score: 0.36521862176312436

Naive Bayes

Model Overview

$$(x = v \mid C_k) = \frac{1}{\sqrt{2\pi\sigma_k^2}} e^{-\frac{(v-\mu_k)^2}{2\sigma_k^2}}$$

Pros:

- Simplicity and speed
- Computational efficiency
- Scalability

Cons:

- Assumes feature independence
- Sensitive
- Often too simple for complex relations

Hyperparameters Tested:

- **'var_smoothing':**
[5e-9, 1e-9, 5e-8, 1e-8, 5e-7, 1e-7, 5e-6, 1e-6, 5e-5, 1e-5]

Best Model:

Best Score: 0.37777

Test Score: 0.36861

MultiLayer Perceptron

Model Overview

Neural network that learns the relationship between linear and non-linear data.

Layers - Input, Hidden, and Output

$$f(x) = f(3)(f(2)(f(1)(x)))$$

Pros:

- Fully Connected Network
- Can handle both structured and unstructured data

Cons:

- Prone to overfitting
- Sensitive to noise data

Hyperparameters Tested:

- **'hidden_layer_sizes': [(50,), (100,), (50,50), (100,50)],**
- **'activation': ['relu', 'tanh'],**
- **'solver': ['adam', 'sgd'],**
- **'alpha': [0.0001, 0.001, 0.01],**

Best Model:

Validation Score: 0.38254

Test Score: 0.36819

Support Vector Machine

Model Overview

Supervised machine learning problem that finds a hyperplane to separate the classes.

$$\forall n: |y_n - (x_n' \beta + b)| \leq \epsilon$$

Pros:

- Memory Efficient
- Can be used for both linear and non-linear classification
- Robust to overfitting

Cons:

- Computational complexity
- Sensitive to noise
- Memory Usage

Hyperparameters Tested:

- **'C':** [0.1, .5, 1, 10],
- **'kernel':** ['linear', 'rbf']
- **'gamma':** ['scale', 'auto']

Best Model:

Validation Score: 0.41289

Test Score: 0.42075

VGG19

Model Overview

VGG-19 is a convolutional neural network that is 19 layers deep.

Pros:

- Great Accuracy
- Simple to implement and understand
- Great at extracting rich features from images

Cons:

- Can be computationally expensive

Hyperparameters Tested:

- **'units'**: [min_value=256, max_value=1024, step=32, (300)],
- **'dropout'**: [0.1, 0.2, 0.3, '0.4', 0.5],
- **'solver'**: ['adam', 'sgd'],
- **'learning_rate'**: [min_value=1e-6, max_value=1e-3 (6.364e-05)],

Best Model:

Validation Score: 0.6549

Test Score: 0.66

ResNet50

Model Overview

ResNet-50 is a 50-layer convolutional neural network (48 convolutional layers, one MaxPool layer, and one average pool layer)

Pros:

- Depth Handling with residuals
- Ability to use pretrained weights

Cons:

- Computationally expensive
- Sensitive to noisy data

Hyperparameters Tested:

- **'units':** [min_value=256, max_value=1024, step=32, (1024)],
- **'dropout':** [0.1, 0.2, 0.3, '0.4', 0.5'],
- **'solver':** ['adam', 'sgd'],
- **'learning_rate':** [min_value=1e-6, max_value=1e-3 (9.8648e-05)],

Best Model:

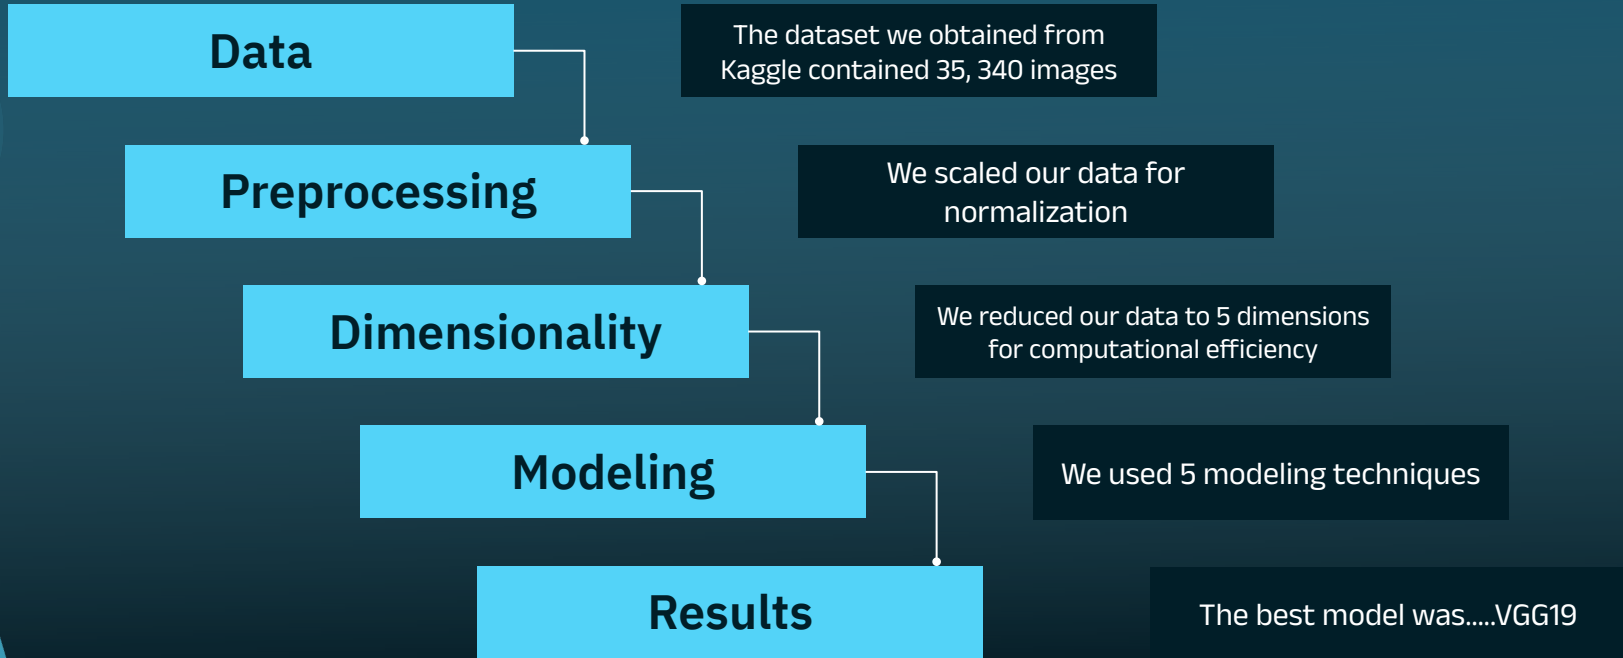
Validation Score: 0.6194

Test Score: 0.62

Results

MODELS	ACCURACY
SGD	36.5%
NB	36.9%
MLP	36.8%
SVM	42%
CNN, VGG19	67%
CNN, ResNet50	62%

Summary and Conclusions





**Mislabeled
Dataset**



**Computational
Resources**

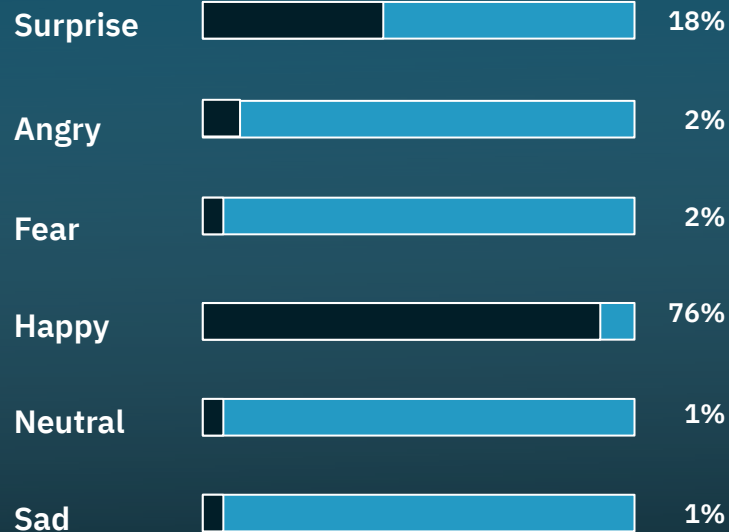
Limitations

**Time
Constraints**

Keras Libraries



Future Applications - Build a Website





Thanks!

Any questions?

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)

References

- [Facial Recognition Dataset \(kaggle.com\)](https://www.kaggle.com/datasets)
- [Image Detection Using the VGG-19 Convolutional Neural Network | by Melisa Bardhi | MLearning.ai | Medium](#)
- [Programming Image Classification with Machine Learning \(kili-technology.com\)](https://kili-technology.com/)
- [Dimensionality Reduction Techniques | Python \(analyticsvidhya.com\)](https://analyticsvidhya.com/)
- [Image classification | TensorFlow Core](#)
- [fit_transform\(\), fit\(\), transform\(\) in Scikit-Learn | Uses & Differences \(analyticsvidhya.com\)](#)
- [Introduction to PCA: Image Compression example | Kaggle](#)
- [Building a Topic Modelling for Images using LDA and Transfer Learning | by Saikumar Jagadeeswaran | Analytics Vidhya | Medium](#)

SCALED


MODELS	TIME	ACCURACY
SGD	3 min 25 sec	29%
MLP	6 mins 7 sec	42%
RF	2 mins 25 sec	47%

SCALED - LDA

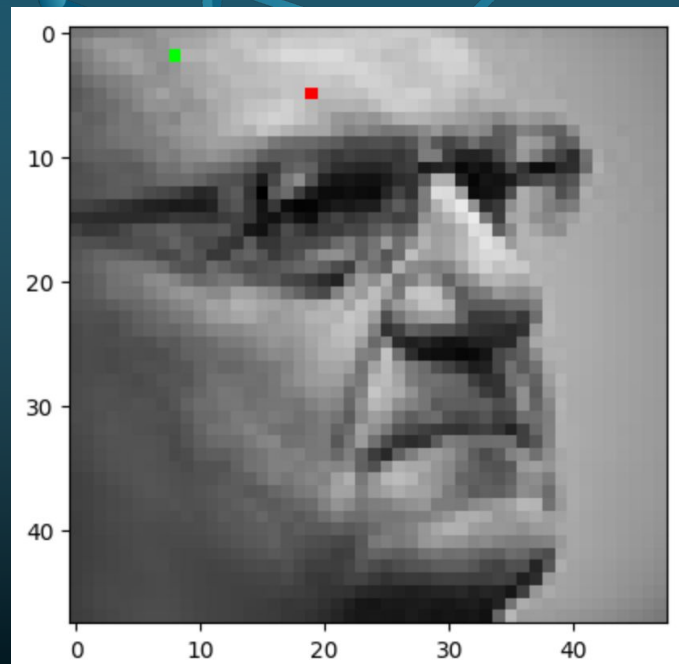
MODELS	TIME	ACCURACY
SGD	445 ms	65%
MLP	21.7 sec	66%
SVM	1 min 48 sec	67%
RF	9.69 sec	64%


```
plt.imshow(X_train[0])  
X_train[0,0,0]
```

```
array([0.46666667, 0.46666667, 0.46666667], dtype=float32)
```



```
X_train[0,5,19] = [1,0,0]  
X_train[0,2,8] = [0,1,0]  
plt.imshow(X_train[0])
```



Limitations & Future Work

- Misabeled Dataset
- Computational Resource Limitations
- Time Constraints
- Keras Libraries (

ORIGINAL

MODELS	TIME	ACCURACY
NB	12 mins	28%
MLP	8 mins 17 sec	37%
RF	4 mins	46%
NB	TIME	ACCURACY
SVM	422 ms	61%
VGG19	18.2 sec	67%
ResNet50	1 min 36 sec	67%
RF	9.82 sec	63%

Random Forest

Model Overview

$$RFf_i = \frac{\sum_{j \in \text{all trees}} normf_{ij}}{T}$$

Pros:

- High accuracy
- Robust to overfitting
- Feature importance
- Parallelization

Cons:

- Model interpretability
- Computational complexity
- Memory usage
- Biased toward dominant classes

Hyperparameters Tested:

- **'n_estimators':** [100, 200, 300]
- **'max_depth':** [None, 5, 10, 20]
- **'min_samples_split':** [2, 5, 10]
- **'min_samples_leaf':** [1, 2, 4]

Best Model:

Validation Score:


Test Score:

The background is a dark blue gradient. It features several light blue geometric elements: lines connecting dots (resembling a network or graph) in the corners, and faint, larger-scale geometric shapes like triangles and circles scattered across the field.


KERAS

Sample Images


Angry



Fear




Happy



Neutral



Sad



Surprise

